

# Generalized Linear Models with brms

Filippo Gambarota 

*filippo.gambarota@unipd.it*

*University of Padova*

# Recap about linear models

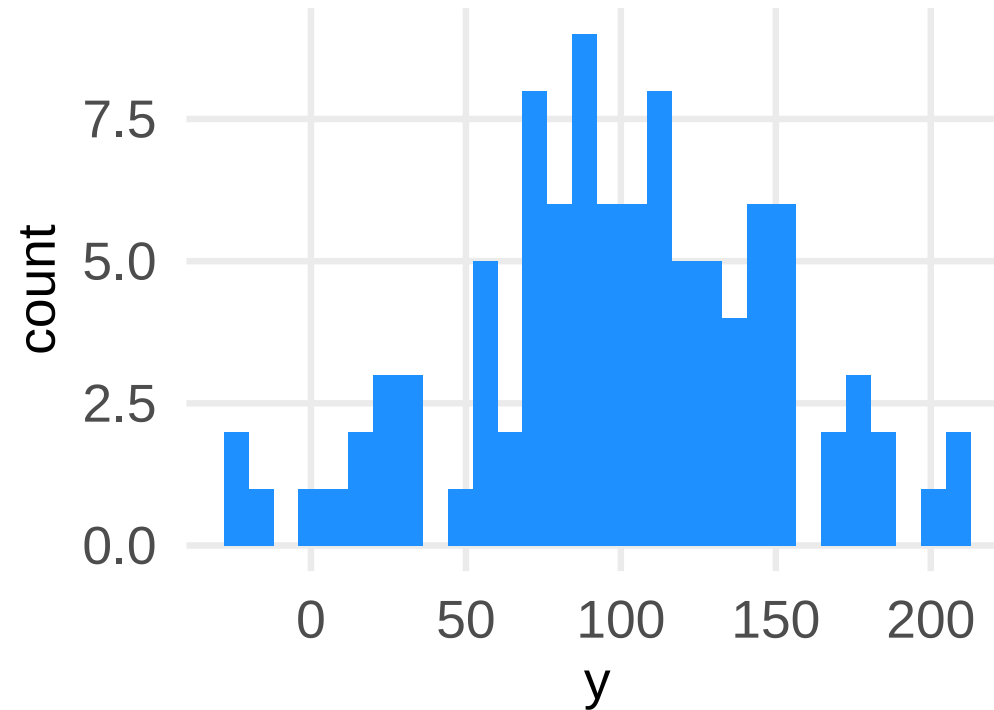
# (almost) everything is a linear model

Most of the statistical analysis that you usually perform, is essentially a linear model.

- ▶ The **t-test** is a linear model where a numerical variable  $y$  is predicted by a factor with two levels  $x$
- ▶ The **one-way anova** is a linear model where a numerical variable  $y$  is predicted by one factor with more than two levels  $x$
- ▶ The **correlation** is a linear model where a numerical variable  $y$  is predicted by another numerical variable  $x$
- ▶ The **ancova** is a linear model where a numerical variable  $y$  is predicted by a numerical variable  $x$  and a factor with two levels  $g$
- ▶ ...

# What is a linear model?

Let's start with a single variable  $y$ . We assume that the variable comes from a Normal distribution:



# What is a linear model?

What we can do with this variable? We can estimate the parameters that define the Normal distribution thus  $\mu$  (the mean) and  $\sigma$  (the standard deviation).

```
mean(y)
#> [1] 100
sd(y)
#> [1] 50
```

Using a linear model we can just fit a model without predictors, also known as intercept-only model.

```
fit <- glm(y ~ 1, family = gaussian(link = "identity"))
summary(fit)
```

# What is a linear model?

```
#>
#> Call:
#> glm(formula = y ~ 1, family = gaussian(link = "identity"))
#>
#> Coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept)      100           5      20  <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> (Dispersion parameter for gaussian family taken to be 2500)
#>
#>      Null deviance: 247500  on 99  degrees of freedom
#> Residual deviance: 247500  on 99  degrees of freedom
#> AIC: 1069.2
```

# What is a linear model?


```
#>  
#> Number of Fisher Scoring iterations: 2
```

I am using `glm` because I want to estimate parameters using Maximul Likelihood, but the results are the same as using `lm`.

Basically we estimated the mean (Intercept) and the standard deviation Dispersion, just take the square root thus 50.

What we are doing is essentially finding the  $\mu$  and  $\sigma$  that maximised the log-likelihood of the model fixing the observed data.


# What is a linear model?



And assuming that we know  $\sigma$  (thus fixing it at 50):




# What is a linear model?




Thus, with the estimates of  $g_{lm}$ , we have this model fitted on the data:

# What is a linear model?



# Including a predictor

When we include a predictor, we are actually try to explain the variability of  $y$  using a variable  $x$ . For example, this is an hypothetical relationship:



# Including a predictor

Seems that there is a positive (linear) relationship between  $x$  and  $y$ . We can try to improve the previous model by adding the predictor:

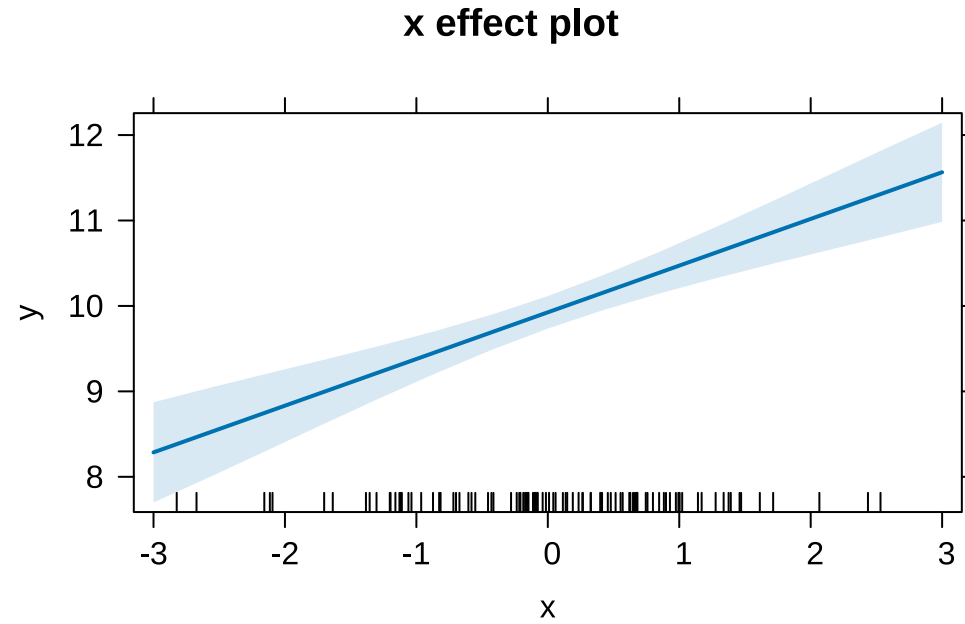
```
fit <- glm(y ~ x, family = gaussian(link = "identity"))
summary(fit)
```

```
#>
#> Call:
#> glm(formula = y ~ x, family = gaussian(link = "identity"))
#>
#> Coefficients:
#>              Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   9.92538    0.09590 103.500  < 2e-16 ***
#> x              0.54648    0.09272   5.894 5.35e-08 ***
#> ---
```

# Including a predictor

```
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> (Dispersion parameter for gaussian family taken to be 0.9193148)
#>
#>      Null deviance: 122.028  on 99  degrees of freedom
#> Residual deviance:  90.093  on 98  degrees of freedom
#> AIC: 279.35
#>
#> Number of Fisher Scoring iterations: 2
```

# Including a predictor



# **Assumptions of the linear model**


# Assumptions of the linear model

include here



# Assumptions of the linear model

More practically, we are saying that the model allows for varying the mean i.e., each  $x$  value can be associated with a different  $\mu$  but with a fixed (and estimated)  $\sigma$ .



# Bayesian Models

# rstanarm

Let's fit the same model but with `rstanarm`. I'm using `rstanarm` just because it's faster, but the idea (and the result) is the same using `brms`.

```
#>
#> SAMPLING FOR MODEL 'continuous' NOW (CHAIN 1).
#> Chain 1:
#> Chain 1: Gradient evaluation took 2.7e-05 seconds
#> Chain 1: 1000 transitions using 10 leapfrog steps per transition would
take 0.27 seconds.
#> Chain 1: Adjust your expectations accordingly!
#> Chain 1:
#> Chain 1:
#> Chain 1: Iteration:      1 / 2000 [  0%]   (Warmup)
#> Chain 1: Iteration:    200 / 2000 [ 10%]   (Warmup)
#> Chain 1: Iteration:    400 / 2000 [ 20%]   (Warmup)
```

# rstanarm

```
#> Chain 1: Iteration: 600 / 2000 [ 30%] (Warmup)
#> Chain 1: Iteration: 800 / 2000 [ 40%] (Warmup)
#> Chain 1: Iteration: 1000 / 2000 [ 50%] (Warmup)
#> Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
#> Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
#> Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
#> Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
#> Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
#> Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
#> Chain 1:
#> Chain 1: Elapsed Time: 0.026 seconds (Warm-up)
#> Chain 1: 0.031 seconds (Sampling)
#> Chain 1: 0.057 seconds (Total)
#> Chain 1:
#>
```

# rstanarm

```
#> SAMPLING FOR MODEL 'continuous' NOW (CHAIN 2).  
#> Chain 2:  
#> Chain 2: Gradient evaluation took 1.2e-05 seconds  
#> Chain 2: 1000 transitions using 10 leapfrog steps per transition would  
take 0.12 seconds.  
#> Chain 2: Adjust your expectations accordingly!  
#> Chain 2:  
#> Chain 2:  
#> Chain 2: Iteration:      1 / 2000 [  0%] (Warmup)  
#> Chain 2: Iteration:    200 / 2000 [ 10%] (Warmup)  
#> Chain 2: Iteration:    400 / 2000 [ 20%] (Warmup)  
#> Chain 2: Iteration:    600 / 2000 [ 30%] (Warmup)  
#> Chain 2: Iteration:    800 / 2000 [ 40%] (Warmup)  
#> Chain 2: Iteration:   1000 / 2000 [ 50%] (Warmup)  
#> Chain 2: Iteration:   1001 / 2000 [ 50%] (Sampling)
```

# rstarm

```
#> Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
#> Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
#> Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
#> Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
#> Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
#> Chain 2:
#> Chain 2: Elapsed Time: 0.026 seconds (Warm-up)
#> Chain 2: 0.033 seconds (Sampling)
#> Chain 2: 0.059 seconds (Total)
#> Chain 2:
#>
#> SAMPLING FOR MODEL 'continuous' NOW (CHAIN 3).
#> Chain 3:
#> Chain 3: Gradient evaluation took 1e-05 seconds
#> Chain 3: 1000 transitions using 10 leapfrog steps per transition would
```

# rstarm

```
take 0.1 seconds.  
#> Chain 3: Adjust your expectations accordingly!  
#> Chain 3:  
#> Chain 3:  
#> Chain 3: Iteration:      1 / 2000 [  0%] (Warmup)  
#> Chain 3: Iteration:    200 / 2000 [ 10%] (Warmup)  
#> Chain 3: Iteration:    400 / 2000 [ 20%] (Warmup)  
#> Chain 3: Iteration:    600 / 2000 [ 30%] (Warmup)  
#> Chain 3: Iteration:    800 / 2000 [ 40%] (Warmup)  
#> Chain 3: Iteration:   1000 / 2000 [ 50%] (Warmup)  
#> Chain 3: Iteration:   1001 / 2000 [ 50%] (Sampling)  
#> Chain 3: Iteration:   1200 / 2000 [ 60%] (Sampling)  
#> Chain 3: Iteration:   1400 / 2000 [ 70%] (Sampling)  
#> Chain 3: Iteration:   1600 / 2000 [ 80%] (Sampling)  
#> Chain 3: Iteration:   1800 / 2000 [ 90%] (Sampling)
```

# rstanarm

```
#> Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
#> Chain 3:
#> Chain 3: Elapsed Time: 0.025 seconds (Warm-up)
#> Chain 3: 0.031 seconds (Sampling)
#> Chain 3: 0.056 seconds (Total)
#> Chain 3:
#>
#> SAMPLING FOR MODEL 'continuous' NOW (CHAIN 4).
#> Chain 4:
#> Chain 4: Gradient evaluation took 1.3e-05 seconds
#> Chain 4: 1000 transitions using 10 leapfrog steps per transition would
take 0.13 seconds.
#> Chain 4: Adjust your expectations accordingly!
#> Chain 4:
#> Chain 4:
```



# rstanarm

```
#> Chain 4: Iteration:      1 / 2000 [  0%] (Warmup)
#> Chain 4: Iteration:    200 / 2000 [ 10%] (Warmup)
#> Chain 4: Iteration:    400 / 2000 [ 20%] (Warmup)
#> Chain 4: Iteration:    600 / 2000 [ 30%] (Warmup)
#> Chain 4: Iteration:    800 / 2000 [ 40%] (Warmup)
#> Chain 4: Iteration:   1000 / 2000 [ 50%] (Warmup)
#> Chain 4: Iteration:   1001 / 2000 [ 50%] (Sampling)
#> Chain 4: Iteration:   1200 / 2000 [ 60%] (Sampling)
#> Chain 4: Iteration:   1400 / 2000 [ 70%] (Sampling)
#> Chain 4: Iteration:   1600 / 2000 [ 80%] (Sampling)
#> Chain 4: Iteration:   1800 / 2000 [ 90%] (Sampling)
#> Chain 4: Iteration:   2000 / 2000 [100%] (Sampling)
#> Chain 4:
#> Chain 4:   Elapsed Time: 0.027 seconds (Warm-up)
#> Chain 4:           0.034 seconds (Sampling)
```

# rstanarm

```
#> Chain 4: 0.061 seconds (Total)
```

```
#> Chain 4:
```

```
#>
```

```
#> Model Info:
```

```
#> function: stan_glm
```

```
#> family: gaussian [identity]
```

```
#> formula: y ~ x
```

```
#> algorithm: sampling
```

```
#> sample: 4000 (posterior sample size)
```

```
#> priors: see help('prior_summary')
```

```
#> observations: 100
```

```
#> predictors: 2
```

```
#>
```

```
#> Estimates:
```

# rstanarm

```
#>           mean    sd   10%   50%   90%
#> (Intercept) 10.0    0.1   9.9  10.0  10.2
#> x           0.3    0.1   0.2   0.3   0.5
#> sigma       1.1    0.1   1.0   1.1   1.2
#>
#> Fit Diagnostics:
#>           mean    sd   10%   50%   90%
#> mean_PPD 10.1    0.2   9.9  10.1  10.3
#>
#> The mean_ppd is the sample average posterior predictive distribution of
the outcome variable (for details see help('summary.stanreg')).
#>
#> MCMC diagnostics
#>           mcse Rhat n_eff
#> (Intercept)  0.0  1.0  3598
```

# rstanarm

```
#> x          0.0  1.0  3420
#> sigma      0.0  1.0  3490
#> mean_PPD    0.0  1.0  3625
#> log-posterior 0.0  1.0  1563
#>
```

#> For each parameter, mcse is Monte Carlo standard error, n\_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence Rhat=1).

```
#> # A draws_df: 1000 iterations, 4 chains, and 3 variables
#>      (Intercept)      x sigma
#> 1      10.1 0.36    1.1
#> 2      10.1 0.31    1.1
#> 3      10.0 0.34    1.1
#> 4      10.2 0.27    1.0
```

# rstanarm

```
#> 5      10.1 0.42    1.0
#> 6       9.9 0.21    1.2
#> 7       9.8 0.27    1.1
#> 8      10.2 0.21    1.2
#> 9      10.2 0.36    1.1
#> 10     9.9 0.28    1.1
#> # ... with 3990 more draws
#> # ... hidden reserved variables {'.chain', '.iteration', '.draw'}
```