

4. Introduzione a Stan

Massimiliano Pastore
Università di Padova


Contents

1 Introduction to STAN

- The Binomial model
 - rstan
 - cmdstanr
- Inference with normal distribution
 - cmdstanr

2 Packages for using STAN

Introduction to STAN



¹Lee, M. D., & Wagenmakers, E. J. (2014). *Bayesian cognitive modeling: A practical course*. Cambridge University Press.

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

STAN modeling language

- A Stan program defines a statistical model through a conditional probability function $p(\theta|y, x)$, where θ is a sequence of modeled unknown values (e.g., model parameters, latent variables, missing data, future predictions), y is a sequence of modeled known values, and x is a sequence of unmodeled predictors and constants (e.g., sizes, hyperparameters).
- Stan programs consist of variable type declarations and statements.

Binomial model: STAN code

```
data {  
  int<lower=0> N;  
  array[N] int<lower=0, upper=1> y;  
}
```

Binomial model: STAN code

```
data {  
  int<lower=0> N;  
  array[N] int<lower=0, upper=1> y;  
}
```

```
parameters {  
  real<lower=0, upper=1> theta;  
}
```

Binomial model: STAN code

```
data {  
  int<lower=0> N;  
  array[N] int<lower=0, upper=1> y;  
}
```

```
parameters {  
  real<lower=0, upper=1> theta;  
}
```

```
model {  
  target += bernoulli_lpmf( y | theta );  
}
```

```
Scode <- "  
  data {  
    int<lower=0> N;  
    array[N] int<lower=0, upper=1> y;  
  }  
  parameters {  
    real<lower=0, upper=1> theta;  
  }  
  model {  
    target += bernoulli_lpmf( y | theta );  
  }  
"
```

rstan

```
# carico il pacchetto  
library(rstan)
```

rstan

```
# carico il pacchetto
library(rstan)
```

```
# lista dei dati
dataList <- list(
  y = c( 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0 ),
  N = 14 )
```

rstan

```
# carico il pacchetto
library(rstan)
```

```
# lista dei dati
dataList <- list(
  y = c( 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0 ),
  N = 14 )
```

```
# fit del modello
fit_binomial.1 <- stan(
  model_code = Scode,
  data = dataList,
  seed = 1,
  cores = parallel::detectCores()
)
```


rstan

```
fit_binomial.1
```

rstan

fit_binomial.1

Inference for Stan model: anon model.

```
4 chains, each with iter=2000; warmup=1000; thin=1;
```

```
post-warmup draws per chain=1000, total post-warmup draws=4000
```

	mean	se_mean	sd	5.5%	94.5%	n_eff	Rhat
theta	0.750	0.003	0.105	0.569	0.898	1463.574	1.000
lp__	-9.512	0.017	0.716	-10.863	-9.000	1810.545	1.001

Samples were drawn using NUTS(diag_e) at Fri May 2 14:34:55

For each parameter, `n_eff` is a crude measure of effective sample size and `Rhat` is the potential scale reduction factor on split chains (indicating convergence, `Rhat=1`).

rstan: traceplot

```
stan_trace( fit_binomial.1, inc_warmup = TRUE )
```

```
stan_trace( fit_binomial.1, inc_warmup = TRUE )
```

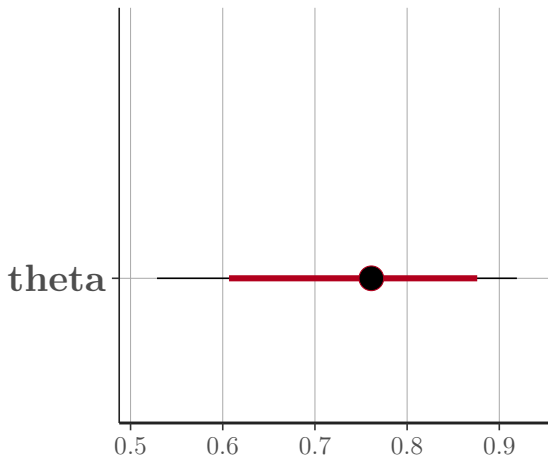


rstan: plotting functions

```
plot( fit_binomial.1 )
```

rstan: plotting functions

```
plot( fit_binomial.1 )
```




rstan: plotting functions

```
plot( fit_binomial.1, plotfun = "hist" )
```

rstan: plotting functions

```
plot( fit_binomial.1, plotfun = "hist" )
```




rstan: plotting functions

```
plot( fit_binomial.1, plotfun = "dens" )
```

rstan: plotting functions

```
plot( fit_binomial.1, plotfun = "dens" )
```



cmdstanr

```
# carico il pacchetto  
library( cmdstanr )
```

cmdstanr

```
# carico il pacchetto  
library( cmdstanr )
```

```
# salvo il codice in un file esterno  
writeLines( Scode, con = "binomial.stan" )
```

cmdstanr

```
# carico il pacchetto  
library( cmdstanr )
```

```
# salvo il codice in un file esterno  
writeLines( Scode, con = "binomial.stan" )
```

```
# compilo il codice  
model <- cmdstan_model( "binomial.stan" )
```

cmdstanr

```
# carico il pacchetto
library( cmdstanr )
```

```
# salvo il codice in un file esterno
writeLines( Scode, con = "binomial.stan" )
```

```
# compilo il codice
model <- cmdstan_model( "binomial.stan" )
```

```
# stima dei parametri
fit_binomial.2 <- model$sample(
  data = dataList,
  seed = 1,
  parallel_chains = parallel::detectCores()
)
```

cmdstanr

```
fit_binomial.2$summary()
```

cmdstanr

```
fit_binomial.2$summary()
```

```
# A tibble: 2 x 10
```


	variable	mean	median	sd	mad	q5	q95	rhat	ess_bulk
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	lp__	-9.51	-9.22	0.705	0.310	-11.0	-9.00	1.00	193
2	theta	0.750	0.759	0.103	0.107	0.574	0.909	1.00	149

cmdstanr: traceplot

```
library( bayesplot )  
draws <- fit_binomial.2$draws()  
mcmc_trace( draws, pars = "theta" )
```

cmdstanr: traceplot

```
library( bayesplot )  
draws <- fit_binomial.2$draws()  
mcmc_trace( draws, pars = "theta" )
```




rstan: plotting functions

```
mcmc_hist( draws, pars = "theta" )
```

rstan: plotting functions

```
mcmc_hist( draws, pars = "theta" )
```




rstan: plotting functions

```
mcmc_dens( draws, pars = "theta" )
```

rstan: plotting functions

```
mcmc_dens( draws, pars = "theta" )
```




rstan: plotting functions

```
mcmc_areas( draws, pars = "theta", prob = 0.8 )
```

rstan: plotting functions

```
mcmc_areas( draws, pars = "theta", prob = 0.8 )
```



Exercise

Proof that the obtained empirical posterior is a $B(1 + 11, 1 + 3)$

Exercise


Proof that the obtained empirical posterior is a $B(1 + 11, 1 + 3)$

```
POST <- extract( fit_binomial.1, 'theta' )  
plot( density( POST$theta ) )  
curve( dbeta( x, 12, 4 ), add = TRUE, col = 'red' )
```

Exercise

Proof that the obtained empirical posterior is a $B(1 + 11, 1 + 3)$

```
POST <- extract( fit_binomial.1, 'theta' )  
plot( density( POST$theta ) )  
curve( dbeta( x, 12, 4 ), add = TRUE, col = 'red' )
```



Inference with normal distribution

Inference with normal distribution

- Perhaps the most useful (or extensively used) probability model for data analysis is the normal distribution:

$$p(y|\mu, \sigma) = \frac{1}{Z} \exp \left(-\frac{1}{2} \frac{(y - \mu)^2}{\sigma^2} \right)$$

where $Z = \sigma\sqrt{2\pi}$.

- The problem of inference involves estimating parameters μ and σ .


Example

- Let's suppose to have assessed the cognitive functions of 10 patients with a test (e.g. QI with average 100 and sd 15), obtaining the following scores:


$$y = \{73, 101, 74, 76, 112, 71, 71, 75, 97, 67\}$$

- The sample mean score is $\bar{y} = 81.7$ (with $s = 15.57$).
- We want to estimate posterior distributions of μ and σ .

Graphical representation of the model




Graphical representation of the model



- Graphical scheme of the model:

- $\mu \sim \text{Normal}(100, 10)$
- $\sigma \sim \text{Uniform}(0, 20)$
- $y_i \sim \text{Normal}(\mu, \sigma)$

Graphical representation of the model



- Graphical scheme of the model:
 - $\mu \sim \text{Normal}(100, 10)$
 - $\sigma \sim \text{Uniform}(0, 20)$
 - $y_i \sim \text{Normal}(\mu, \sigma)$
- The value assigned to the maximum of variance (in this case 20) must be chosen according to the problem.
- In this case the wide spread of σ value indicates a high level of uncertainty about this parameter.

Example

```
data {  
  int N; // sample size  
  vector[N] y; // data  
}
```

Example

```
data {  
  int N; // sample size  
  vector[N] y; // data  
}  
parameters {  
  real mu;  
  real<lower=0, upper=20> sigma;  
}
```

Example

```
data {  
  int N; // sample size  
  vector[N] y; // data  
}  
parameters {  
  real mu;  
  real<lower=0, upper=20> sigma;  
}  
model {  
  target += normal_lpdf( mu | 100, 10 ); // mu prior  
  target += normal_lpdf( y | mu, sigma ); // likelihood  
}
```

Example

```
data {  
  int N; // sample size  
  vector[N] y; // data  
}  
parameters {  
  real mu;  
  real<lower=0, upper=20> sigma;  
}  
model {  
  target += normal_lpdf( mu | 100, 10 ); // mu prior  
  target += normal_lpdf( y | mu, sigma ); // likelihood  
}  
generated quantities {  
  real ypred;  
  ypred = normal_rng(mu, sigma);  
}
```

```
normal_stancode <- "  
  data {  
    int N; // sample size  
    vector[N] y; // data  
  }  
  parameters {  
    real mu;  
    real<lower=0, upper=20> sigma;  
  }  
  model {  
    target += normal_lpdf( mu | 100, 10 ); // mu prior  
    target += normal_lpdf( y | mu, sigma ); // likelihood  
  }  
  generated quantities {  
    real ypred;  
    ypred = normal_rng(mu, sigma);  
  }  
"
```

rstan

```
dataList <- list( y = c( 73, 101, 74, 76, 112, 71, 71,  
                        75, 97, 67 ), N = 10 )
```

rstan

```
dataList <- list( y = c( 73, 101, 74, 76, 112, 71, 71,  
                        75, 97, 67 ), N = 10 )
```

```
fit_normal.1 <- stan(  
  model_code = normal_stancode,  
  data = dataList,  
  seed = 1,  
  cores = parallel::detectCores()  
)
```


rstan

```
print( fit_normal.1, pars = c("mu","sigma") )
```

rstan

```
print( fit_normal.1, pars = c("mu","sigma") )
```

Inference for Stan model: anon_model.

4 chains, each with iter=2000; warmup=1000; thin=1;

post-warmup draws per chain=1000, total post-warmup draws=4000

	mean	se_mean	sd	5.5%	94.5%	n_eff	Rhat
mu	85.48	0.11	4.70	78.16	93.21	1701.20	1
sigma	15.79	0.07	2.45	11.72	19.51	1387.05	1

Samples were drawn using NUTS(diag_e) at Fri May 2 14:41:29


For each parameter, n_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (should be close to 1).
convergence, Rhat=1).

rstan:traceplot

```
stan_trace( fit_normal.1, pars = c("mu","sigma"),  
            inc_warmup = TRUE )
```

rstan:traceplot

```
stan_trace( fit_normal.1, pars = c("mu","sigma"),  
            inc_warmup = TRUE )
```




rstan: plotting functions

```
plot( fit_normal.1, plotfun = "hist",  
      pars = c("mu","sigma") )
```

rstan: plotting functions

```
plot( fit_normal.1, plotfun = "hist",  
      pars = c("mu","sigma") )
```





rstan: posterior predictive check

```
y <- dataList$y
yrep <- matrix(
  extract( fit_normal.1, pars = "ypred" )$ypred,
  ncol = length( y ) )
ppc_dens_overlay( y, yrep[1:50,])
```

rstan: posterior predictive check

```
y <- dataList$y
yrep <- matrix(
  extract( fit_normal.1, pars = "ypred" )$ypred,
  ncol = length( y ) )
ppc_dens_overlay( y, yrep[1:50,])
```





⁴Figure is produced by the `stat_density2d()` function (ggplot2 package).

100%

```
# salvo il codice in un file esterno
writeLines( normal_stancode, con = "normal.stan" )
```

100

```
# salvo il codice in un file esterno
writeLines( normal_stancode, con = "normal.stan" )
```

```
# compilo il codice
model_normal <- cmdstan_model( "normal.stan" )
```

cmdstanr

```
# salvo il codice in un file esterno
writeLines( normal_stancode, con = "normal.stan" )
```

```
# compilo il codice
model_normal <- cmdstan_model( "normal.stan" )
```

```
# stima dei parametri
fit_normal.2 <- model_normal$sample(
  data = dataList,
  seed = 1,
  parallel_chains = parallel::detectCores()
)
```

cmdstanr


```
fit_normal.2$summary()
```


cmdstanr: traceplot

```
draws <- fit_normal.2$draws()
mcmc_trace( draws, pars = c("mu","sigma") )
```

cmdstanr: traceplot

```
draws <- fit_normal.2$draws()
mcmc_trace( draws, pars = c("mu","sigma") )
```




rstan: plotting functions

```
mcmc_hist( draws, pars = c("mu","sigma") )
```

rstan: plotting functions

```
mcmc_hist( draws, pars = c("mu","sigma") )
```




rstan: plotting functions

```
mcmc_dens( draws, pars = c("mu","sigma") )
```

rstan: plotting functions

```
mcmc_dens( draws, pars = c("mu","sigma") )
```



rstan: plotting functions

```
mcmc_areas( draws, pars = c("mu","sigma"), prob = 0.8 )
```

```
mcmc_areas( draws, pars = c("mu","sigma"), prob = 0.8 )
```




cmdstanr: posterior predictive check

```
yrep <- fit_normal.2$draws("ypred", format = "matrix")  
yrep <- matrix( yrep, ncol = length( y ) )  
ppc_dens_overlay( y, yrep[1:50,] )
```

cmdstanr: posterior predictive check

```
yrep <- fit_normal.2$draws("ypred", format = "matrix")  
yrep <- matrix( yrep, ncol = length( y ) )  
ppc_dens_overlay( y, yrep[1:50,] )
```



Packages for using STAN

- **rstan**: It provides the R interface to Stan. The **rstan** package allows one to conveniently fit Stan models from R and access the output, including posterior inferences and intermediate quantities such as evaluations of the log posterior density and its gradients.
- **cmdstanr**: a lightweight interface to Stan for R users that provides an alternative to the traditional **rstan** interface.

Packages for using STAN

- **rethinking**: This package accompanies a book and course on Bayesian data analysis (by Richard McElreath), featured Quadratic approximate estimation through `quap()` and Hamiltonian Monte Carlo through `ulam()`.

Example with rethinking: Quadratic approximate

```
library( rethinking )
fit_quap <- quap(
  alist(
    y ~ dnorm( mu, sigma ),
    mu ~ dnorm( 100, 10 ),
    sigma ~ dunif( 0, 20 )
  ), data = dataList, start = list( mu = 90, sigma = 10))
```

Example with rethinking: Quadratic approximate

```
library( rethinking )
fit_quap <- quap(
  alist(
    y ~ dnorm( mu, sigma ),
    mu ~ dnorm( 100, 10 ),
    sigma ~ dunif( 0, 20 )
  ), data = dataList, start = list( mu = 90, sigma = 10))

precis( fit_quap )
```

Example with rethinking: Quadratic approximate

```
library( rethinking )
fit_quap <- quap(
  alist(
    y ~ dnorm( mu, sigma ),
    mu ~ dnorm( 100, 10 ),
    sigma ~ dunif( 0, 20 )
  ), data = dataList, start = list( mu = 90, sigma = 10))
```

```
precis( fit_quap )
```

	mean	sd	5.5%	94.5%
mu	85.12	4.52	77.90	92.33
sigma	15.17	3.54	9.51	20.83

Example with rethinking

```
library( rethinking )
fit <- ulam(
  alist(
    y ~ dnorm( mu, sigma ),
    mu ~ dnorm( 100, 10 ),
    sigma ~ dunif( 0, 20 )
  ),
  data = dataList, chains = 4, cores = 4 )
```

Example with rethinking

```
library( rethinking )  
fit <- ulam(  
  alist(  
    y ~ dnorm( mu, sigma ),  
    mu ~ dnorm( 100, 10 ),  
    sigma ~ dunif( 0, 20 )  
  ),  
  data = dataList, chains = 4, cores = 4 )  
  
precis( fit )
```

Example with rethinking

```
library( rethinking )
fit <- ulam(
  alist(
    y ~ dnorm( mu, sigma ),
    mu ~ dnorm( 100, 10 ),
    sigma ~ dunif( 0, 20 )
  ),
  data = dataList, chains = 4, cores = 4 )
```

```
precis( fit )
```

	mean	sd	5.5%	94.5%	n_eff	Rhat4
mu	85.41	4.29	78.79	92.6	764	1
sigma	15.46	2.31	11.69	19.1	876	1

Packages for using STAN

- **rstanarm**: The goal of this package is to make Bayesian estimation routine for the most common regression models that applied researchers use. This will enable researchers to avoid the counter-intuitiveness of the frequentist approach to probability and statistics with only minimal changes to their existing Rscripts.

Example with rstanarm

```
library( rstanarm )  
d <- data.frame( y = dataList$y )  
stan_glm( y ~ 1, data = d )
```

Example with rstanarm

```
library( rstanarm )
d <- data.frame( y = dataList$y )
stan_glm( y ~ 1, data = d )
```

```
stan_glm
family:      gaussian [identity]
formula:     y ~ 1
observations: 10
predictors:  1
```

```
-----
                Median MAD_SD
(Intercept) 81.7      5.1
```

Auxiliary parameter(s):

```
                Median MAD_SD
sigma 16.2      3.7
```

Packages for using STAN

- **brms**: Fit Bayesian generalized (non-)linear multivariate multilevel models using Stan for full Bayesian inference. A wide range of distributions and link functions are supported, allowing users to fit – among others – linear, robust linear, count data, survival, response times, ordinal, zero-inflated, hurdle, and even self-defined mixture models all in a multilevel context.

Example with brms

```
library( brms )  
brm( y ~ 1, data = d )
```

Example with brms

```
library( brms )
brm( y ~ 1, data = d )
```

```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: y ~ 1
Data: d (Number of observations: 10)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	78.78	3.96	71.27	87.02	1.00	2513	2396

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	15.15	3.40	10.04	23.28	1.00	2280	2003

Packages using STAN

- **shinystan**: graphical user interface.
- **blavaan**: Fit a variety of Bayesian latent variable models, including confirmatory factor analysis, structural equation models, and latent growth curve models.
- **edstan**: Stan for item response theory, it attempts to make easy the fitting of standard item response theory models using rstan.
- ...

For further information

BUGS:



<http://bayesmodels.com/>



[http://www.mrc-bsu.cam.ac.uk/software/bugs/
the-bugs-project-bugs-resources-online/](http://www.mrc-bsu.cam.ac.uk/software/bugs/the-bugs-project-bugs-resources-online/)

rjags:



[http://www.johnmyleswhite.com/notebook/2010/08/20/
using-jags-in-r-with-the-rjags-package/](http://www.johnmyleswhite.com/notebook/2010/08/20/using-jags-in-r-with-the-rjags-package/)

STAN:



<http://mc-stan.org/>



[https://www.mzes.uni-mannheim.de/socialsciencedatalab/
article/applied-bayesian-statistics/](https://www.mzes.uni-mannheim.de/socialsciencedatalab/article/applied-bayesian-statistics/)

Used R packages

- `bayesplot`. Gabry J, Mahr T (2025).
- `brms`. Bürkner P (2017).
- `cmdstanr`. Gabry J, Češnovar R, Johnson A, Bronder S (2025).
- `cowplot`. Wilke C (2024).
- `ggplot2`. Wickham H (2016).
- `knitr`. Xie Y (2025).
- `posterior`. Bürkner P, Gabry J, Kay M, Vehtari A (2025).
- R. R Core Team (2025).
- `Rcpp`. Eddelbuettel D, Francois R, Allaire J, Ushey K, Kou Q, Russell N, Ucar I, Bates D, Chambers J (2025).
- `report`. Makowski D, Lüdtke D, Patil I, Thériault R, Ben-Shachar M, Wiernik B (2023).
- `rethinking`. McElreath R (2024).
- `rstan`. Stan Development Team (2025).
- `rstanarm`. Goodrich B, Gabry J, Ali I, Brilleman S (2024).
- `StanHeaders`. Stan Development Team (2020).



massimiliano.pastore@unipd.it
<https://psicostat.dpss.psy.unipd.it/>



L^AT_EX

